

## Question 1:

### Introduction

In this problem, I implemented the K-nearest neighbors (KNN) algorithm to recognize handwritten digits using the MNIST dataset. The dataset is made of 60,000 training images and 10,000 test images. Each image is represented as a 28x28 pixel grid. The goal was to classify the test images by comparing them with training images using KNN with different distance metrics.

```
import math
import numpy as np
from download_mnist import load
import operator
import time

# classify using kNN
#x_train = np.load('../x_train.npy')
#y_train = np.load('../y_train.npy')
#x_test = np.load('../x_test.npy')
#y_test = np.load('../y_test.npy')
x_train, y_train, x_test, y_test = load()
x_train = x_train.reshape(60000,28,28)  #60,000 images for training (28x28
pixels)
x_test = x_test.reshape(10000,28,28)    #10,000 images for testing (28x28
pixels)
x_train = x_train.astype(float)          #converts pixel value to float
x_test = x_test.astype(float)            #converts pixel value to float
```

I first loaded the MNIST dataset using the load() function from download\_mnist.py, which I modified slightly because the original code wasn't working. I simply changed it to download the MNIST dataset from another website. Since the images were originally in a flattened format (784-dimensional vectors), I reshaped them into 28x28 matrices for better manipulation. I only used a subset of data to improve computation efficiency, as training on the entire dataset would take a long time. I played around with the number of training samples and test samples and ended up with 40,000 training samples and 1,000 test samples, which gave me 95% accuracy when testing it using the L2 distance metric.

```

# newInput: test image to classify
# dataSet: training dataset
# labels: corresponds to labels for dataSet
# k: number of nearest neighbors to consider
def kNNClassify(newInput, dataSet, labels, k, metric):
    result=[]
    #####
    # Input your code here #
    #####
    for test_sample in newInput:
        if metric == 'L1':
            distances = np.sum(np.abs(dataSet - test_sample), axis=(1, 2))
# L1 (Manhattan) Distance
            elif metric == 'L2': # Default is L2 (Euclidean)
                distances = np.sqrt(np.sum((dataSet - test_sample) ** 2,
axis=(1, 2))) # L2 (Euclidean) Distance

            # get indices of k nearest neighbors
            k_neighbors = np.argsort(distances)[:k]

            k_labels = labels[k_neighbors] # get corresponding labels

            # predict the most common label among the k neighbors
            unique_labels, counts = np.unique(k_labels, return_counts=True)
            result.append(unique_labels[np.argmax(counts)])

    #####
    # End of your code #
    #####
    return np.array(result)

```

I implemented the KNN algorithm in the kNNClassify function. The function calculates the distances between a given test image and all training images using either the L1 distance (Manhattan distance) or the L2 distance (Euclidean distance). L1 computes the sum of absolute differences between corresponding pixels, and L2 computes the square root of the sum of squared differences between corresponding pixels. Once the distances were computed, I identified the k nearest neighbors by sorting the distances and selecting the smallest k values. I then determined the most common label among these neighbors and assigned it as the predicted class for the test image.

## Experimentation:

```
num_train_samples = 40000
num_test_samples = 1000
k_value = 6

# Run KNN with L1 distance
start_time_L1 = time.time()
output_labels_L1 = kNNClassify(x_test[:num_test_samples],
x_train[:num_train_samples], y_train[:num_train_samples], k_value,
metric='L1')
accuracy_L1 = (1 - np.count_nonzero(y_test[:num_test_samples] -
output_labels_L1) / len(output_labels_L1))
time_L1 = time.time() - start_time_L1

# Run KNN with L2 distance
start_time_L2 = time.time()
output_labels_L2 = kNNClassify(x_test[:num_test_samples],
x_train[:num_train_samples], y_train[:num_train_samples], k_value,
metric='L2')
accuracy_L2 = (1 - np.count_nonzero(y_test[:num_test_samples] -
output_labels_L2) / len(output_labels_L2))
time_L2 = time.time() - start_time_L2

# Print results
print(f"--- KNN Classification Results on MNIST ---")
print(f"L1 (Manhattan) Distance -> Accuracy: {accuracy_L1:.4f}, Execution
Time: {time_L1:.4f} sec")
print(f"L2 (Euclidean) Distance -> Accuracy: {accuracy_L2:.4f}, Execution
Time: {time_L2:.4f} sec")
```

I ran the KNN classification using many different parameters. The parameters I changed were the number of the training dataset, the number of the testing dataset, and the value of k. Ultimately, to achieve an accuracy of 95% or greater for using the L2 distance metric, I came down to a training size of 40,000, a test size of 1,000, and a k value of 6. As a result, I got:

```
--- KNN Classification Results on MNIST ---
L1 (Manhattan) Distance -> Accuracy: 0.9420, Execution Time: 171.4629 sec
L2 (Euclidean) Distance -> Accuracy: 0.9540, Execution Time: 150.5221 sec
```

From the results, I observed that the L2 distance metric provided better accuracy compared to the L1 distance. This is expected because the Euclidean distance better preserves the spatial structure of images. The final accuracy using L2 distance was over 95%, meeting the requirement of the task. In my case, I noticed a lower execution time for L2 compared to L1. However, logically, it would make sense for L2 to be slightly higher due to additional square root computations.

## Question 2: Linear Classifier

### Introduction:

In this problem, I implemented a linear classifier to recognize handwritten digits using the MNIST dataset, similar to question 1. Instead of using a more complex neural network, I trained a simple linear model with weights  $W$  and bias  $b$ , optimizing it using cross-entropy loss and gradient descent. I also experimented with random search to observe different weight initializations and evaluate their impact on the accuracy.

```
import numpy as np
import pickle
from download_mnist import load

x_train, y_train, x_test, y_test = load()

x_train = x_train / 255.0
x_test = x_test / 255.0

# one-hot encoding function
def one_hot_encode(labels, num_classes):
    one_hot = np.zeros((labels.size, num_classes))
    one_hot[np.arange(labels.size), labels] = 1
    return one_hot

# one-hot encode the labels
num_classes = 10
y_train_onehot = one_hot_encode(y_train, num_classes)
y_test_onehot = one_hot_encode(y_test, num_classes)
```

Similar to question 1, I first loaded the MNIST dataset using the `load()` function. Since pixel values in MNIST range from 0 to 255, I normalized them to the  $[0, 1]$  range by dividing by 255.0. Since the classifier needs to predict one of ten (0-9), I converted the labels into a one-hot encoded format. One-hot encoding represents each label as a 10-dimensional vector where only the correct class index is set to 1, and the rest are 0. This is important for cross-entropy loss. I initialized the model with Random weights  $W$  with small values to prevent large initial gradients, and bias  $b$  was set to zero to keep it simple.

```
num_features = 28 * 28 # 784 for 28x28 pixel images

np.random.seed(42) # initialize weights and bias
W = np.random.randn(num_features, num_classes) * 0.01
b = np.zeros((1, num_classes))
```

```
def cross_entropy_loss(y_true, y_pred): # cross-entropy loss function
    m = y_true.shape[0]
    loss = -np.sum(y_true * np.log(y_pred + 1e-9)) / m
    return loss

def softmax(logits): # softmax function
    exp_scores = np.exp(logits - np.max(logits, axis=1, keepdims=True))
    return exp_scores / np.sum(exp_scores, axis=1, keepdims=True)
```

Since this is a multi-class classification problem, I used the softmax function to convert the model's raw output into a probability distribution across the 10 classes. For training, I used cross-entropy loss.

```
def train_linear_classifier(x_train, y_train_onehot, learning_rate=0.1,
epochs=100): # training function
    global W, b
    for epoch in range(epochs):
        # compute logits and predictions
        logits = np.dot(x_train, W) + b
        y_pred = softmax(logits)

        # compute loss
        loss = cross_entropy_loss(y_train_onehot, y_pred)

        # gradient computation
        m = x_train.shape[0]
        dW = np.dot(x_train.T, (y_pred - y_train_onehot)) / m
        db = np.sum(y_pred - y_train_onehot, axis=0, keepdims=True) / m

        # update weights and bias
        W -= learning_rate * dW
        b -= learning_rate * db

        if (epoch + 1) % 10 == 0:
            print(f"Epoch {epoch + 1}/{epochs}, Loss: {loss:.4f}")
```

To update the model, I performed gradient descent to minimize the loss function. I ran 100 epochs with a learning rate of 0.1 and printed the loss every 10 epochs to track progress.

```
def compute_accuracy(x, y_true_onehot): # Evaluate accuracy
    logits = np.dot(x, W) + b
    y_pred = np.argmax(softmax(logits), axis=1)
    y_true = np.argmax(y_true_onehot, axis=1)
    return np.mean(y_pred == y_true)

train_linear_classifier(x_train, y_train_onehot) # train the linear
classifier

accuracy = compute_accuracy(x_test, y_test_onehot) # evaluate on the
testing set
print(f"Test accuracy: {accuracy * 100:.2f}%")
```

After training, I tested the model's accuracy on the test set. I computed predictions using Softmax and selected the class with the highest probability as the final predicted label. The test accuracy was 87.06%, as shown below. This makes sense because it's a very basic linear model without deep learning enhancements.

```
Epoch 10/100, Loss: 1.6152
Epoch 20/100, Loss: 1.2221
Epoch 30/100, Loss: 1.0147
Epoch 40/100, Loss: 0.8897
Epoch 50/100, Loss: 0.8063
Epoch 60/100, Loss: 0.7464
Epoch 70/100, Loss: 0.7011
Epoch 80/100, Loss: 0.6654
Epoch 90/100, Loss: 0.6365
Epoch 100/100, Loss: 0.6126
Test accuracy: 87.06%
```

```

def random_search(x_train, y_train_onehot, num_searches=50):
    best_accuracy = 0
    best_W = None
    best_b = None

    # poop code (not really sure about this part)
    for i in range(num_searches):
        # generate random weights and biases
        random_W = np.random.randn(num_features, num_classes) * 0.01
        random_b = np.random.randn(1, num_classes) * 0.01

        # compute accuracy for the current random parameters
        logits = np.dot(x_test, random_W) + random_b
        y_pred = np.argmax(softmax(logits), axis=1)
        y_true = np.argmax(y_test_onehot, axis=1)
        accuracy = np.mean(y_pred == y_true)

        if accuracy > best_accuracy:
            best_accuracy = accuracy
            best_W = random_W
            best_b = random_b

    print(f"Best accuracy after random search: {best_accuracy *
100:.2f}%")
    return best_W, best_b

# perform random search
best_W, best_b = random_search(x_train, y_train_onehot)

```

As per directions, I also implemented a random search. Instead of training, I randomly generated 50 different sets of  $W$  and  $b$ , evaluated their accuracy, and kept track of the best-performing parameters. The best accuracy I found using random search was 16.44%, showing that random search is not really good compared to gradient-based optimization which makes sense because we are randomly guessing the parameters.

```
Best accuracy after random search: 16.44%
```

### Conclusion:

To conclude, a linear classifier is simple and easy to implement, but its accuracy is limited because it only learns linear decision boundaries. Using gradient descent with cross-entropy loss ensures proper optimization, but the model is not powerful enough to achieve high



accuracy on the MNIST dataset. Finally, a random search showed that weight initialization affects accuracy, but training is still necessary for optimal performance.